

Game Architecture

Audio

Today's Agenda

- What is sound and how do we represent it?
- Attributes of sound in 3D
- Sound engine architecture
- Ad hoc methods


“The secret to great graphics
is great audio.”

What is sound?

- Delta in local atmospheric pressure
 - Waves oscillate in propagation direction
- Constant tone is periodic sound
 - Frequency (Hz) is $1 / \text{Period}$
- Pitch is $1 / \text{Longest_Period}$
 - Fundamental harmonic
- Loudness is root mean square of pressure
 - Decibels are units of loudness on log scale

What is sound?

- **Delta in local atmospheric pressure**
 - **Waves oscillate in propagation direction**
- Constant tone is periodic sound
 - Frequency (Hz) is $1 / \text{Period}$
- Pitch is $1 / \text{Longest_Period}$
 - Fundamental harmonic
- Loudness is root mean square of pressure
 - Decibels are units of loudness on log scale



What is sound?

- Delta in local atmospheric pressure
 - Waves oscillate in propagation direction
- **Constant tone is periodic sound**
 - **Frequency (Hz) is $1 / \text{Period}$**
- Pitch is $1 / \text{Longest_Period}$
 - Fundamental harmonic
- Loudness is root mean square of pressure
 - Decibels are units of loudness on log scale

Typical adult can hear:


- From 20 Hz up to 20 kHz
- Most sensitive from 2 kHz to 5 kHz
 - Perceptually loudest

Dogs: 67 Hz - 45 kHz

Some bats: Up to 200 kHz!?

What is sound?

- Delta in local atmospheric pressure
 - Waves oscillate in propagation direction
- Constant tone is periodic sound
 - Frequency (Hz) is $1 / \text{Period}$
- **Pitch is $1 / \text{Longest_Period}$**
 - **Fundamental harmonic**
- Loudness is root mean square of pressure
 - Decibels are units of loudness on log scale



What is sound?

- Delta in local atmospheric pressure
 - Waves oscillate in propagation direction
- Constant tone is periodic sound
 - Frequency (Hz) is 1 / Period
- Pitch is 1 / Longest_Period
 - Fundamental harmonic
- **Loudness is root mean square of pressure**
 - **Decibels are units of loudness on log scale**

$$x_{\text{rms}} = \sqrt{\frac{1}{n}(x_1^2 + x_2^2 + \dots + x_n^2)}$$

$$\text{SPL} = 20 \log_{10} \frac{\overset{\text{Pressure} \rightarrow P}{\text{Standard pressure level}}}{\underset{\text{Standard reference pressure} \rightarrow P_0}{P_0}} \text{ dB}$$

Representing sound

- Linear time-invariant system
- Convolution of time shifted impulses
- Samples
- Formats

Representing sound



- **Linear time-invariant system**
 - Convolution of time shifted impulses
 - Samples
 - Formats
- Linear
 - Response to N stimuli is the sum of the responses which would have been caused by each stimulus individually
 - If A produces response X and input B produces response Y then input $(A + B)$ produces response $(X + Y)$
 - *(Wikipedia)*
 - Time-invariant
 - Time shifting input shifts output equally

Representing sound

- Linear time-invariant system
- **Convolution of time shifted impulses**
- Samples
- Formats



Ideal sampling:

① A cosine





② Its spectrum

③ Halving the period...





④ ...doubles the frequency

⑤ A general signal





⑥ Typical spectrum

⑦ Impulse train



⑧

⑨ Sampled signal



⑩ $\omega_c < \frac{1}{2} \omega_s$

Representing sound

- Linear time-invariant system
 - Convolution of time shifted impulses
 - **Samples**
 - Formats
- Sampling rate
 - 8 kHz - Telephone
 - 32 kHz - Wii hardware output
 - 44.1 kHz - CD standard
 - 48 kHz - Modern standard output
 - Nyquist-Shannon sampling theorem
 - Use rate double highest frequency for lossless sampling
 - Bits per sample
 - 16 - Typical input/output
 - 32 - Typical intermediate
 - Channel count
 - 1 - mono
 - 2 - stereo
 - N - Surround, language tracks, other?

Representing sound

- Linear time-invariant system
 - Convolution of time shifted impulses
 - Samples
 - **Formats**
- PCM
 - Pulse code modulation
 - $\text{Size} = \text{Rate} * \text{Bits} * \text{Channels}$
 - Compute = Minimum
 - ADPCM
 - Adaptive differential pulse code modulation
 - Adaptively storing deltas between samples
 - Typically fixes bits_per_sample = 4
 - Compute = Small
 - MP3, Ogg Vorbis, and friends
 - Variable bit rate, lossy
 - Compute = Relatively expensive

Attributes of Sound in 3D


- Distance attenuation
- Doppler
- Panning
- Reverb
- Models of ears and head

Attributes of Sound in 3D

- **Distance attenuation**
 - Doppler
 - Panning
 - Reverb
 - Models of ears and head
- Linear falloff: $1 / \text{distance}$
 - Drop to zero at some max distance

Attributes of Sound in 3D

- Distance attenuation
- **Doppler**
- Panning
- Reverb
- Models of ears and head



$$f' = \left(\frac{v + v'}{v - v_s} \right) f_s$$

Attributes of Sound in 3D

- Distance attenuation
- Doppler
- **Panning**
- Reverb
- Models of ears and head

Constant power pan law:

- $A_{\text{left}} = A * \sin(\pi/2 * \text{pan})$
- $A_{\text{right}} = A * \cos(\pi/2 * \text{pan})$

The “3 dB rule”:


- If sound is coming out of 2 speakers equally, reduce by 3 dB.

Attributes of Sound in 3D

- Distance attenuation
- Doppler
- Panning
- **Reverb**
- Models of ears and head

Components of a sound:

- Dry - sound directly from emitter
- Wet - sound reflected in environment
 - This is reverb.



Sound designers pad their offices to eliminate wet sound.

Attributes of Sound in 3D

- Distance attenuation
- Doppler
- Panning
- Reverb
- **Models of ears and head**

Brain knows a lot about environment from:

- Delta b/w left and right ears
- Delta b/w direct sound and first reflection

We may want to model subtle effects of:

- Ears (their shape and their number)
- Distortion caused by sound traveling through skull (head related transfer function, or HRTF)

But then again, player's speaker config probably has a greater impact on the experience...

Sound engine architecture: Example API

```
// The system
snd_system snd_system_create();
void snd_system_destroy(snd_system system);
void snd_system_set_listener(vec3 pos, quat orient);

// Banks
snd_bank snd_bank_load(snd_system system, string file_name);
void snd_bank_unload(snd_bank bank);

// Groups
snd_group snd_group_create(snd_system system, string name, snd_group parent);
void snd_group_destroy(snd_group group);
void snd_group_set_volume(snd_group group, float volume);
void snd_group_pause(snd_group group, bool paused);

// Instances
snd_instance_handle snd_play(snd_group group, snd_cue cue);
void snd_stop(snd_instance_handle sound);
```

Sound engine architecture: Example API

Owns audio resources and pipeline.

```
// The system
snd_system snd_system_create();
void snd_system_destroy(snd_system system);
void snd_system_set_listener(vec3 pos, quat orient);

// Banks
snd_bank snd_bank_load(snd_system system, string file_name);
void snd_bank_unload(snd_bank bank);

// Groups
snd_group snd_group_create(snd_system system, string name, snd_group parent);
void snd_group_destroy(snd_group group);
void snd_group_set_volume(snd_group group, float volume);
void snd_group_pause(snd_group group, bool paused);

// Instances
snd_instance_handle snd_play(snd_group group, snd_cue cue);
void snd_stop(snd_instance_handle sound);
```

Sound engine architecture: Example API

```
// The system
snd_system snd_system create();
void ;
void t orient);

// Banks
snd_bank snd_bank_load(snd_system system, string file_name);
void snd_bank_unload(snd_bank bank);

// Groups
snd_group snd_group_create(snd_system system, string name, snd_group parent);
void snd_group_destroy(snd_group group);
void snd_group_set_volume(snd_group group, float volume);
void snd_group_pause(snd_group group, bool paused);

// Instances
snd_instance_handle snd_play(snd_group group, snd_cue cue);
void snd_stop(snd_instance_handle sound);
```

A collection of sound clips (compressed PCM files) managed as a unit.

Sound engine architecture: Example API

```
// The system
snd_system snd_system_create();
void snd_system_destroy(snd_system system);
void snd_system_set_listener(vec3 pos, quat orient);

// Banks
snd_bank snd_bank_load(snd_system system, string file_name);
```

Hierarchical, high level mixing control.

```
// Groups
snd_group snd_group_create(snd_system system, string name, snd_group parent);
void snd_group_destroy(snd_group group);
void snd_group_set_volume(snd_group group, float volume);
void snd_group_pause(snd_group group, bool paused);

// Instances
snd_instance_handle snd_play(snd_group group, snd_cue cue);
void snd_stop(snd_instance_handle sound);
```


Sound engine architecture: Example API

```
// The system
snd_system snd_system_create();
void snd_system_destroy(snd_system system);
void snd_system_load(snd_system system, const char *file_name);

// Background
snd_background snd_background_create(snd_system system, const char *file_name);
void snd_background_play(snd_background background, snd_cue cue);

// Groups
snd_group snd_group_create(snd_system system, const char *name, snd_group parent);
void snd_group_play(snd_group group, snd_cue cue);
void snd_group_set_volume(snd_group group, float volume);
void snd_group_destroy(snd_group group);

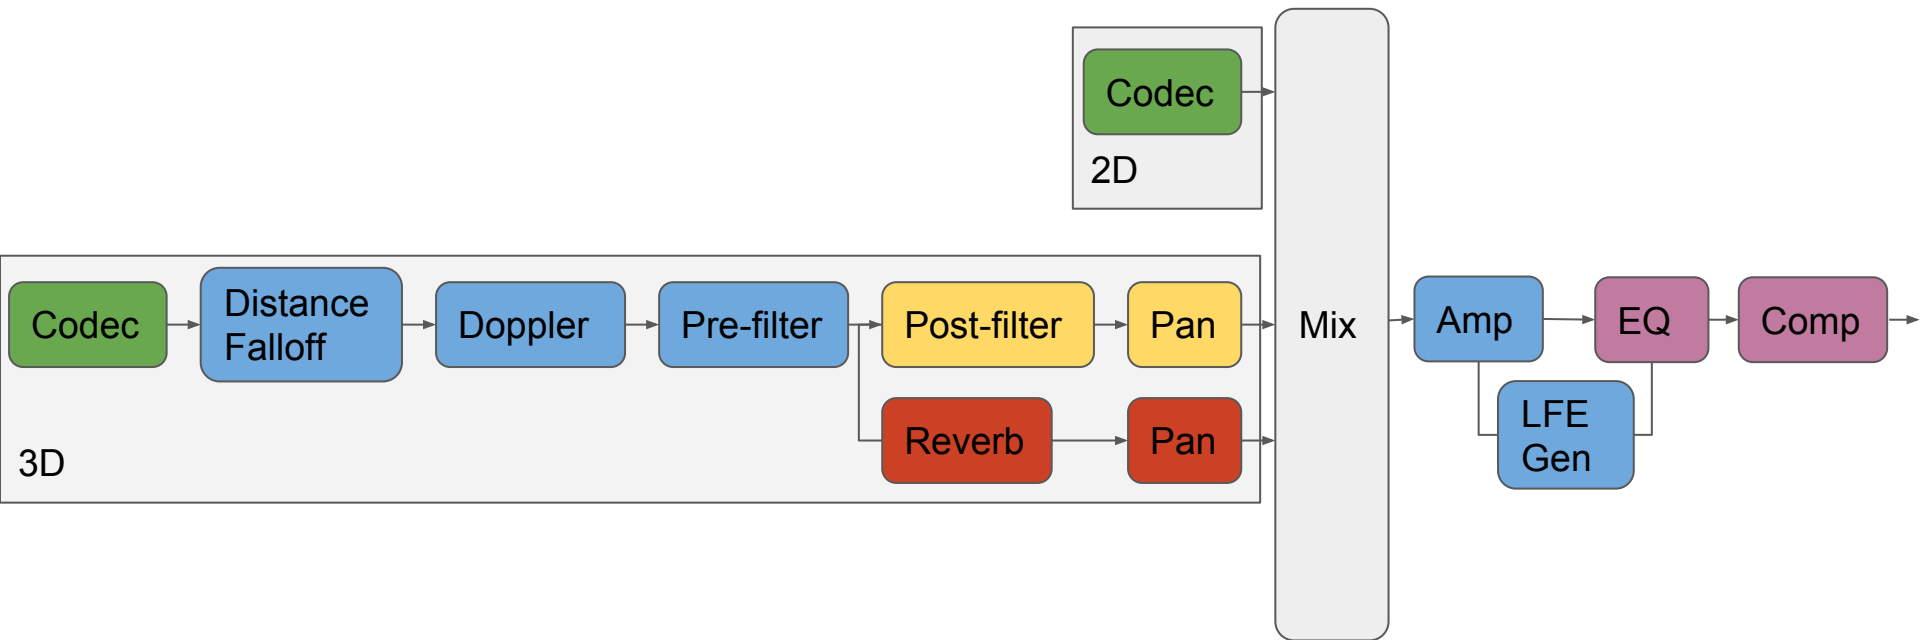
// Instances
snd_instance_handle snd_play(snd_group group, snd_cue cue);
void snd_stop(snd_instance_handle sound);
```

Sound cues are collections of sound clips with some metadata (loop points, mixing information, clip selection criteria).

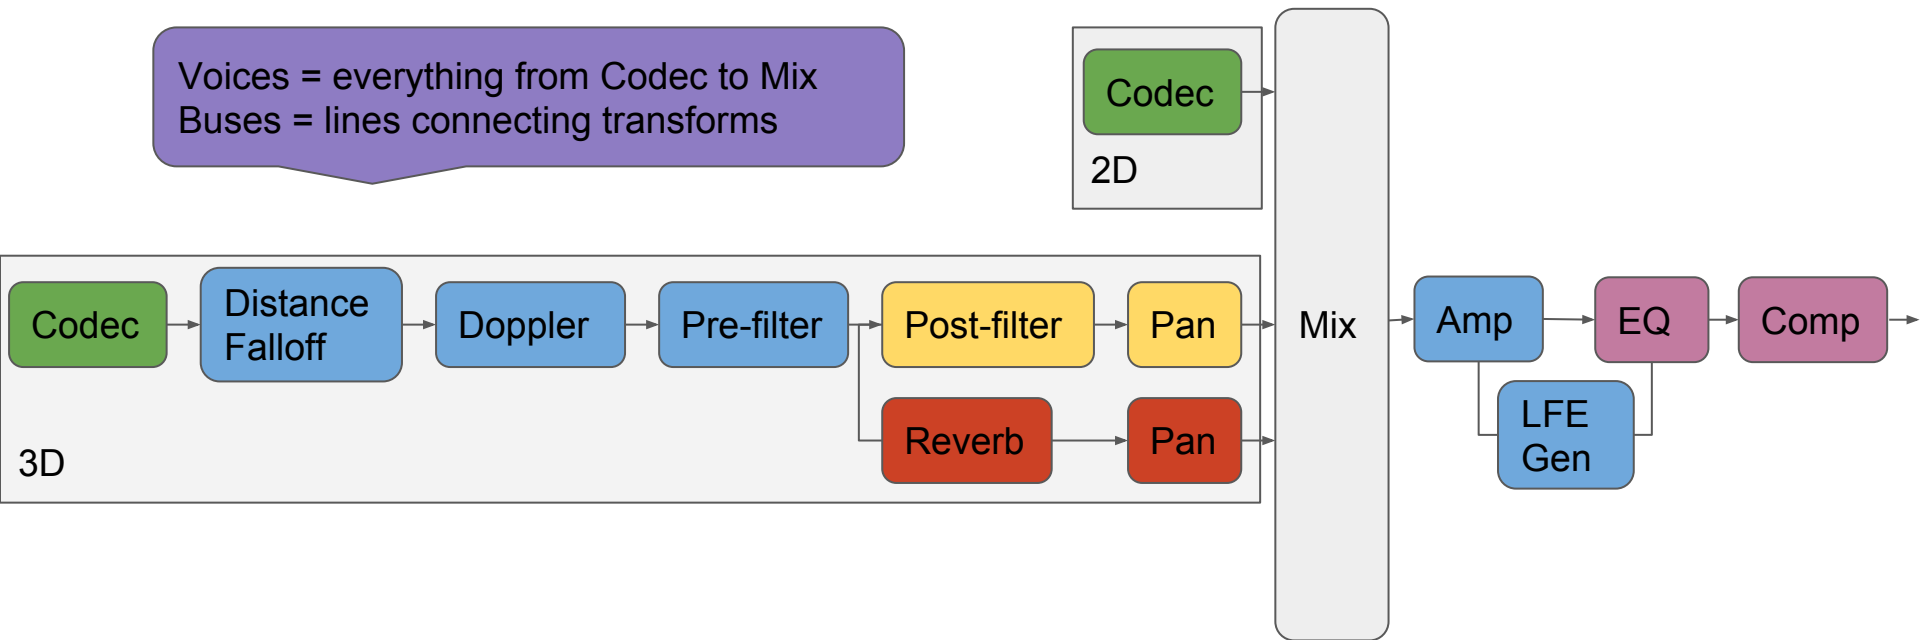
Playing a sound cue allocates a virtual voice and selects a clip.

Sound engine activates virtual voices in range of the listener up to some cap. Active voices go through the mixing pipeline.

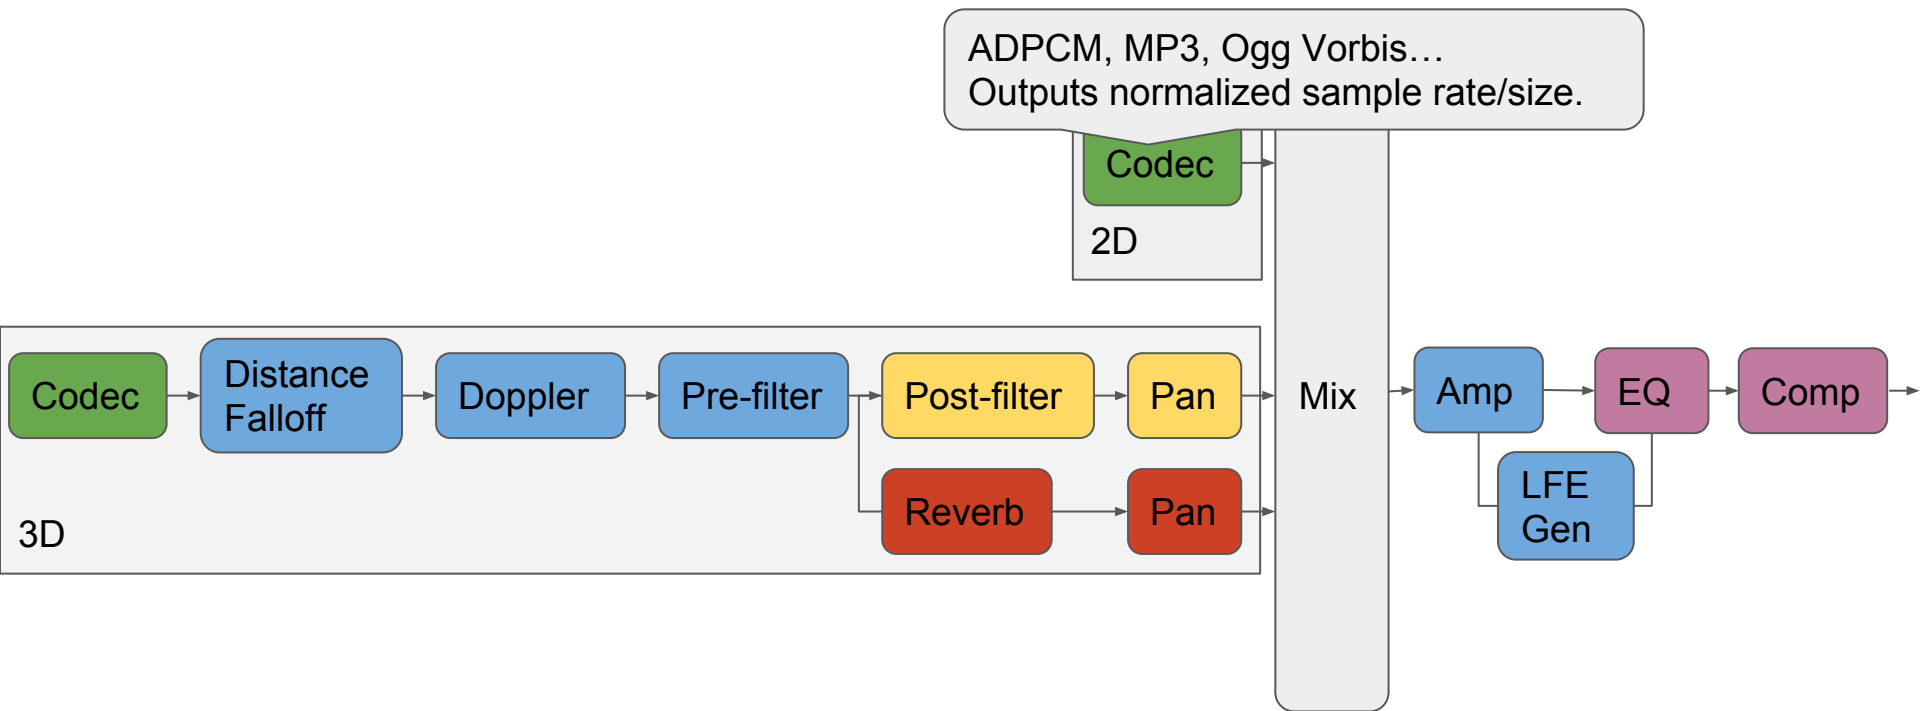
Sound engine architecture: Example Pipeline



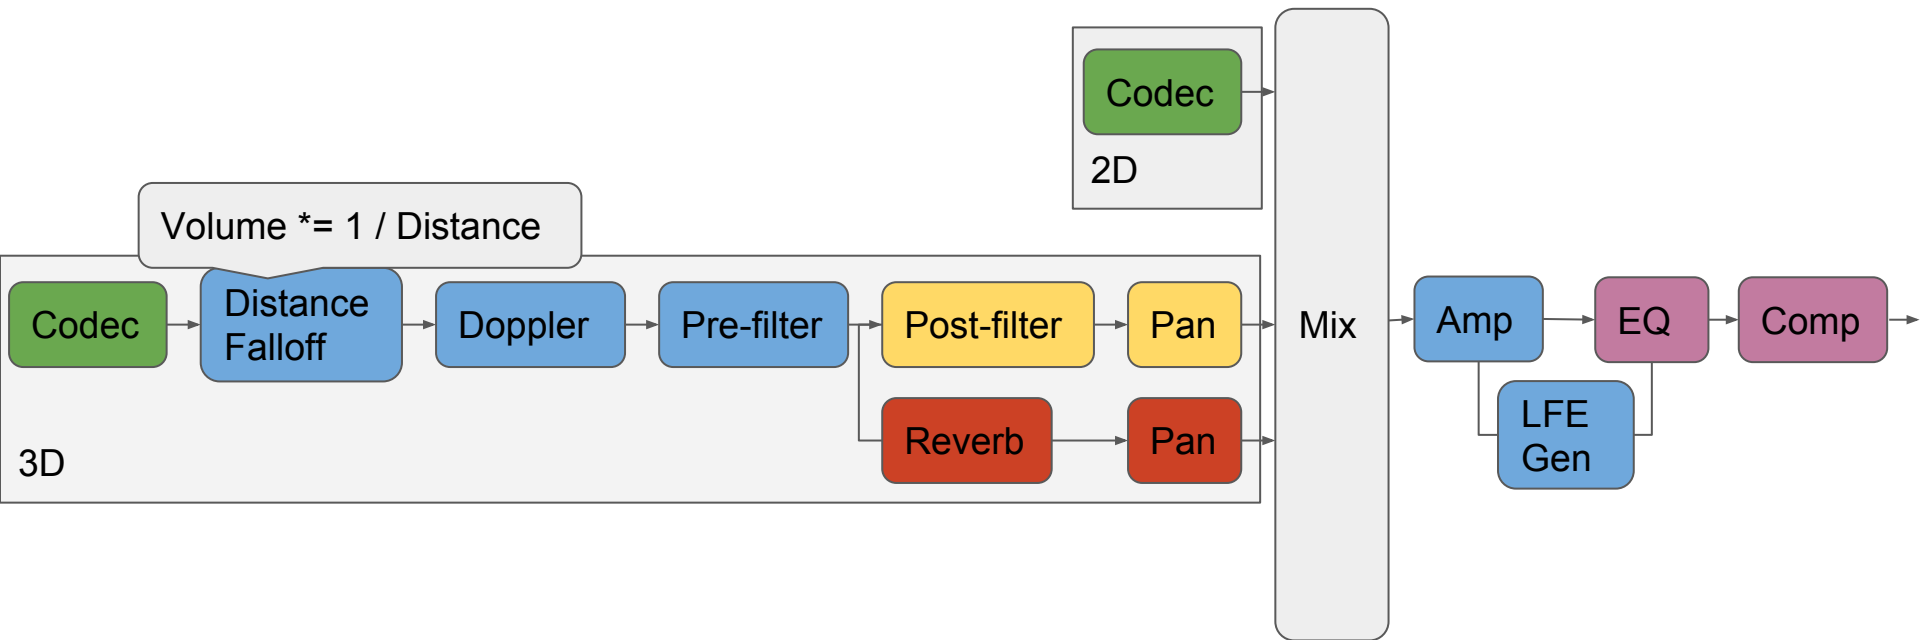
Sound engine architecture: Example Pipeline



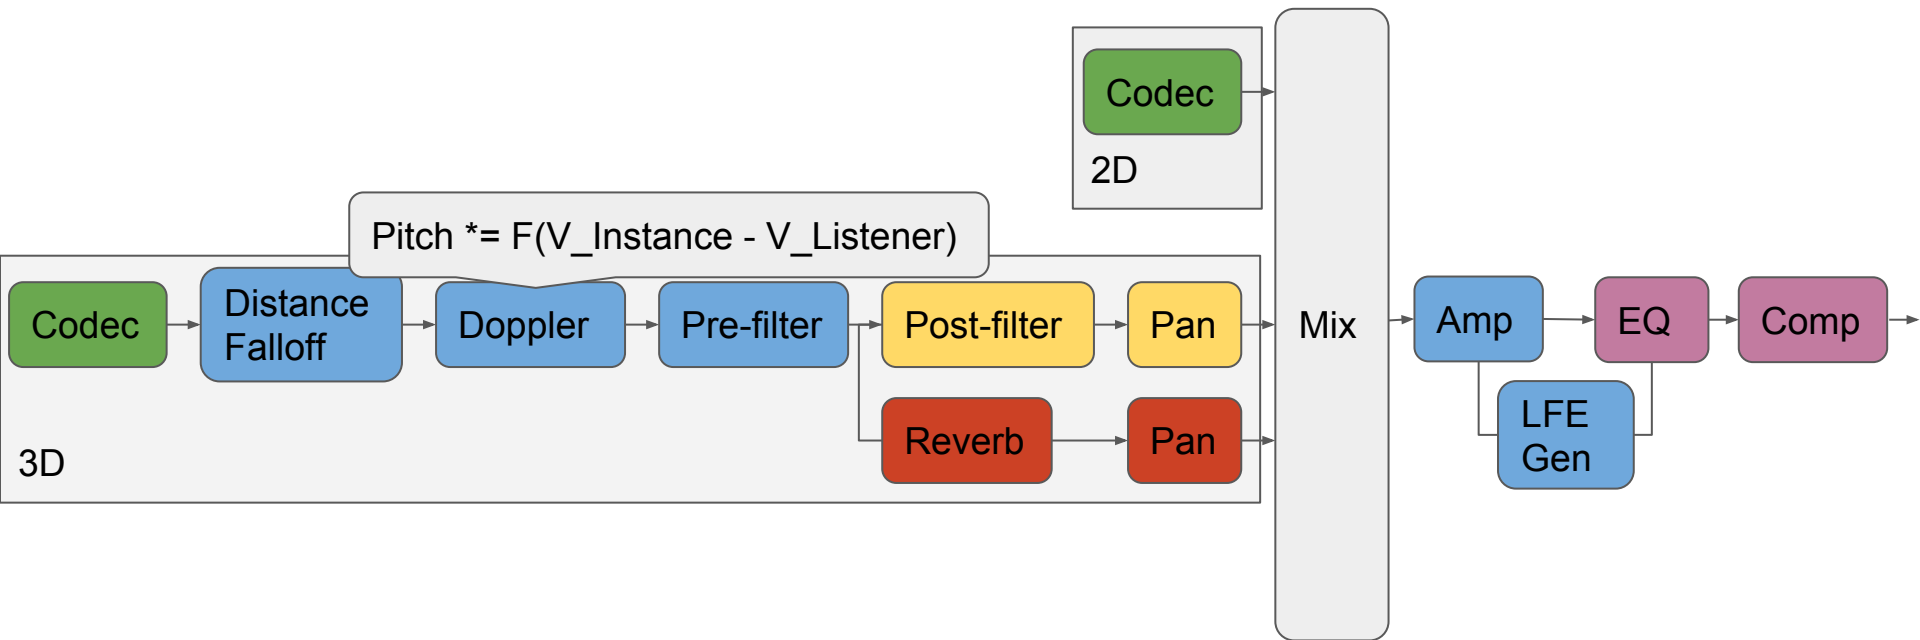
Sound engine architecture: Example Pipeline



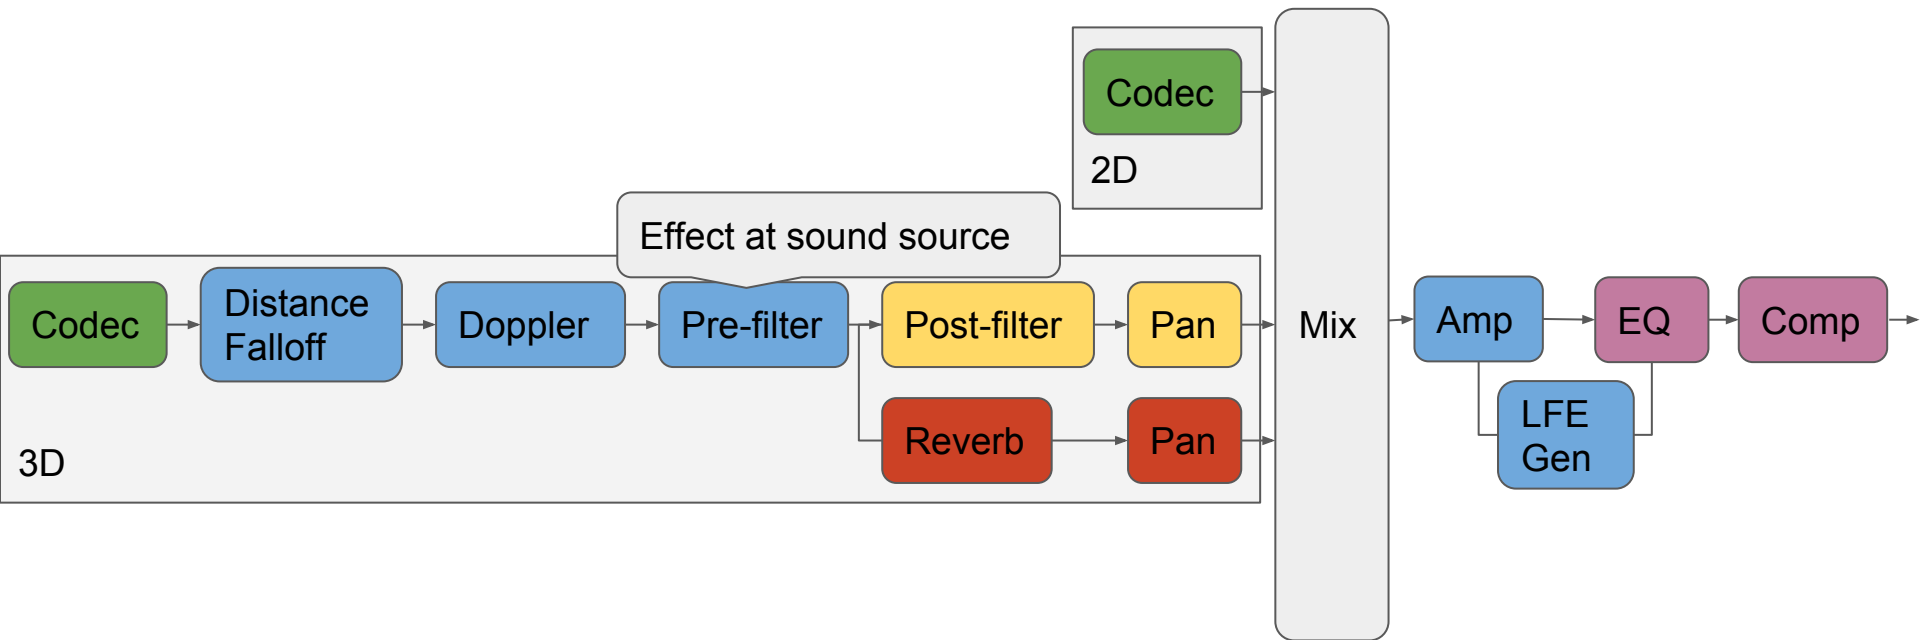
Sound engine architecture: Example Pipeline



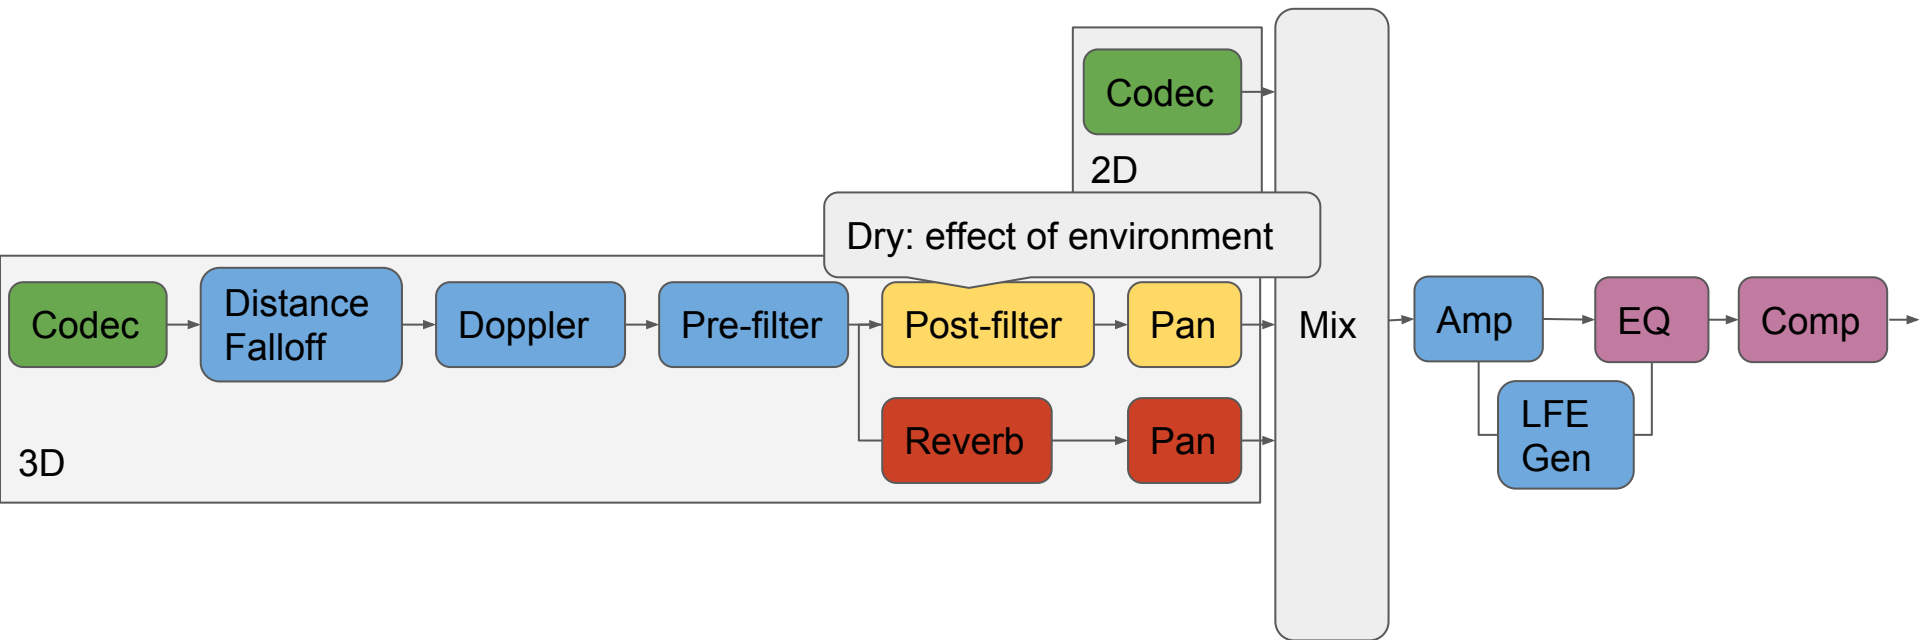
Sound engine architecture: Example Pipeline



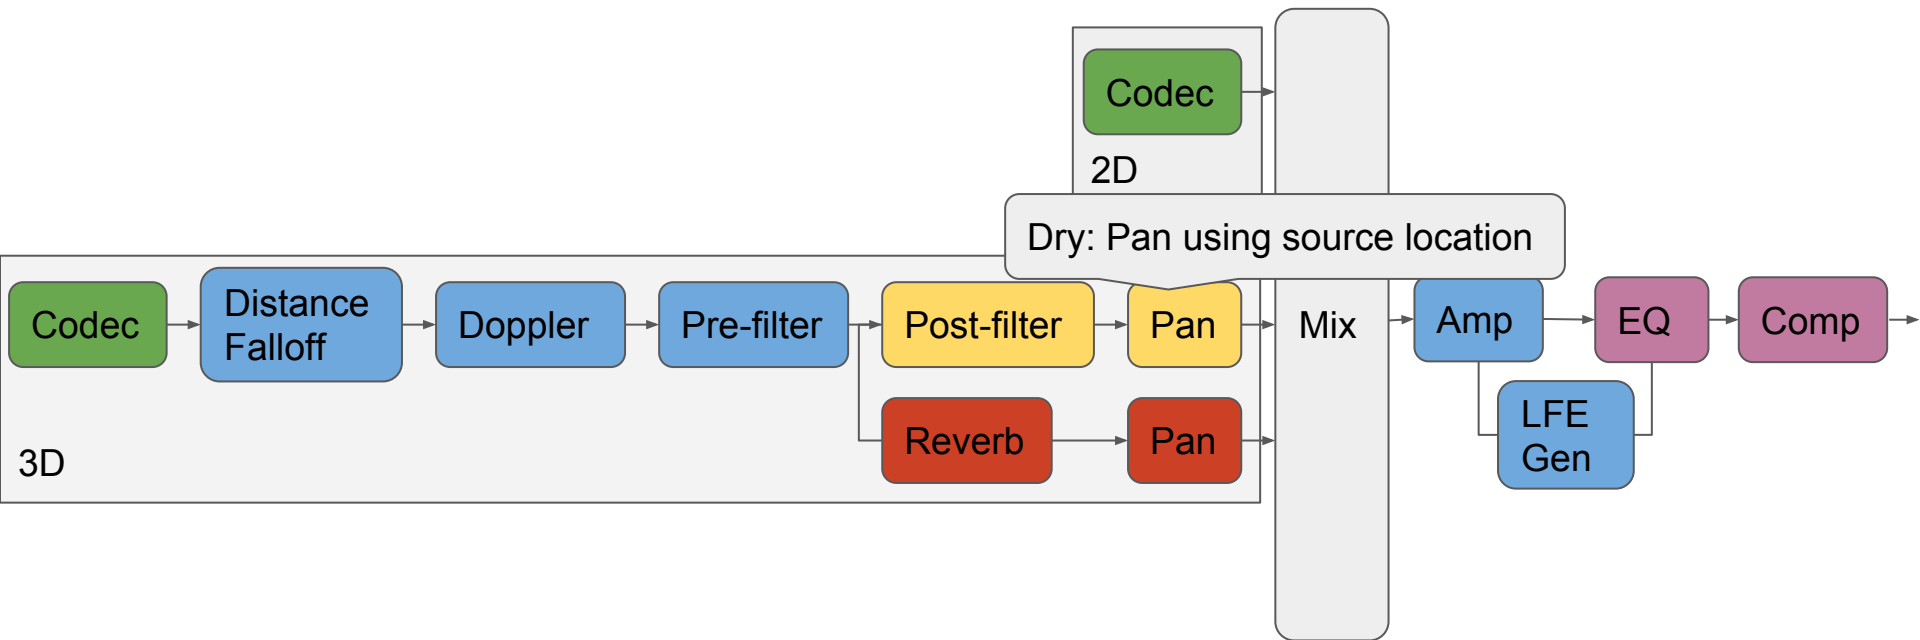
Sound engine architecture: Example Pipeline



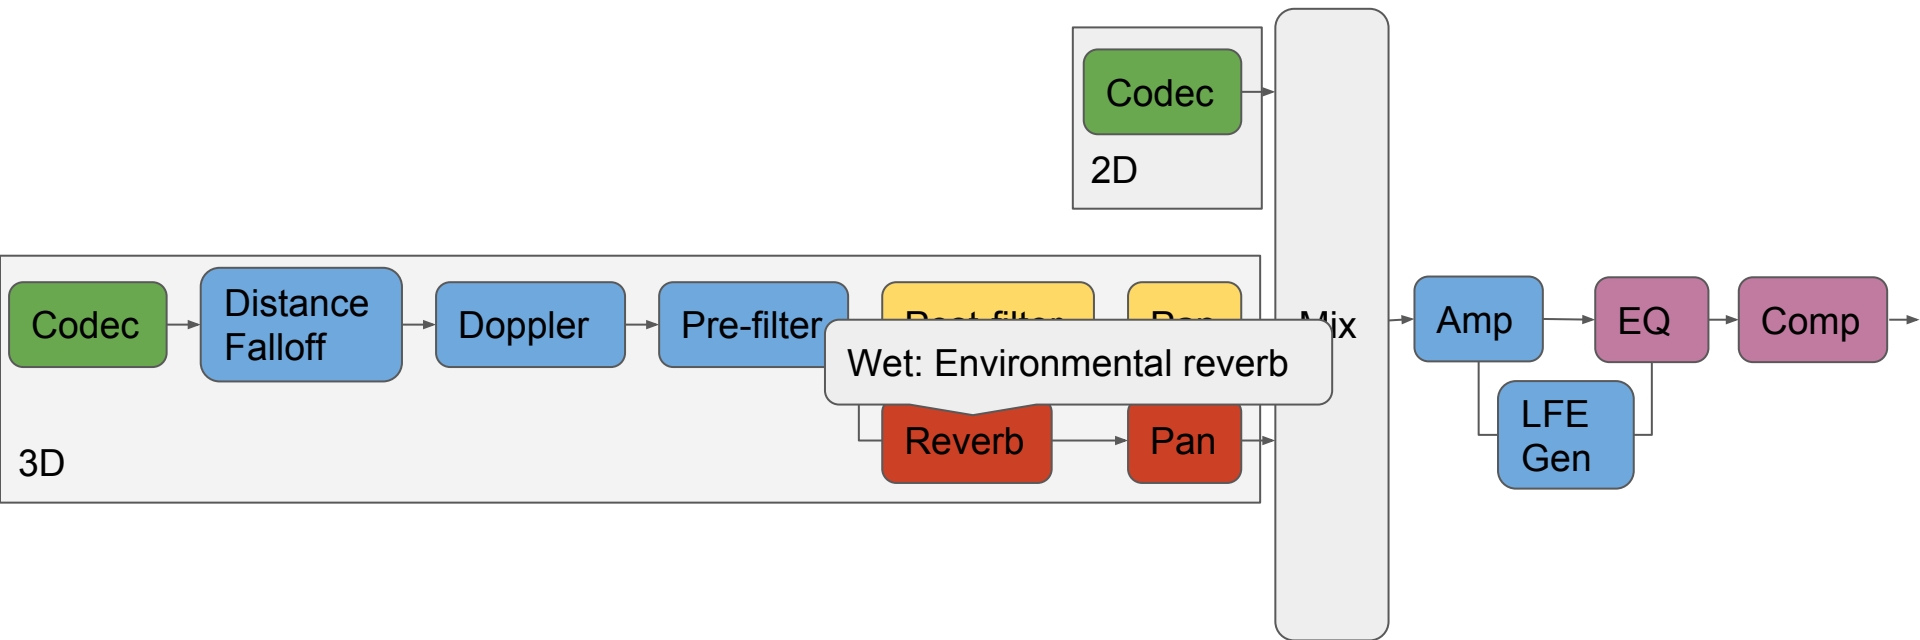
Sound engine architecture: Example Pipeline



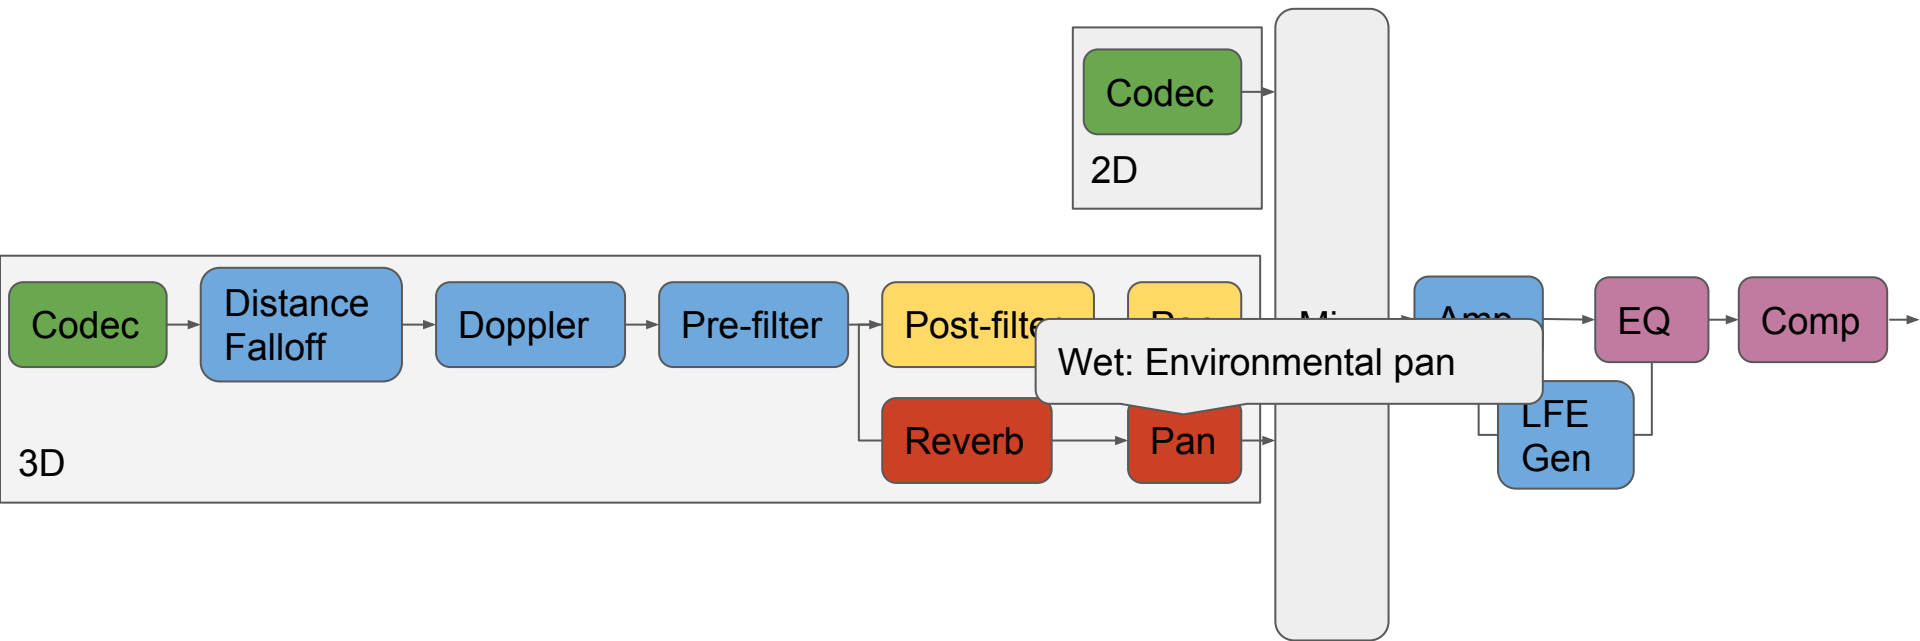
Sound engine architecture: Example Pipeline



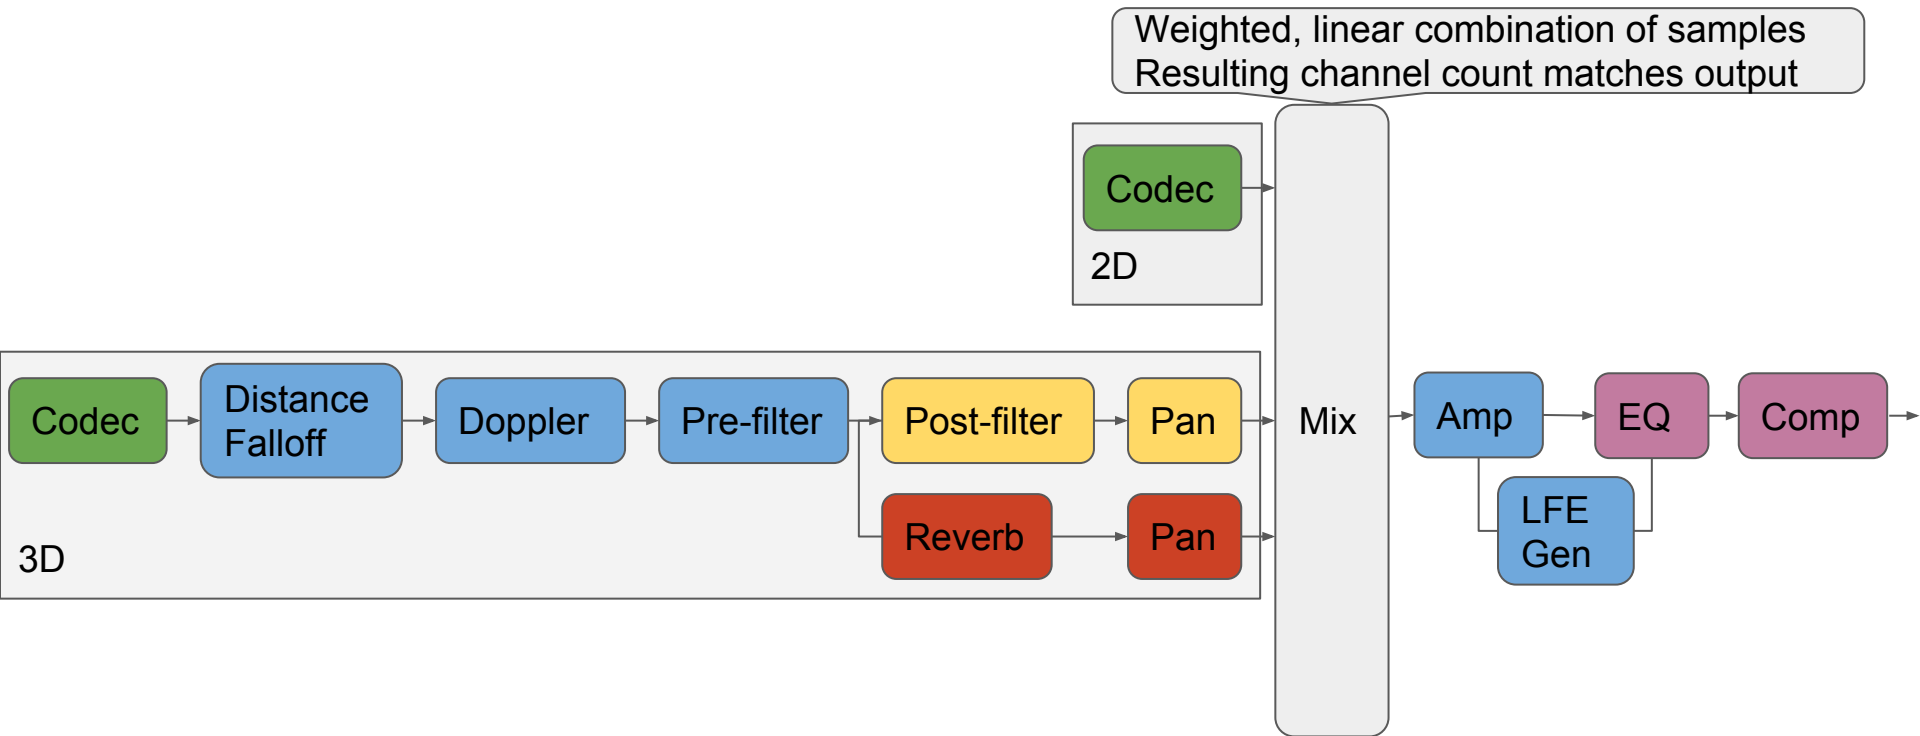
Sound engine architecture: Example Pipeline



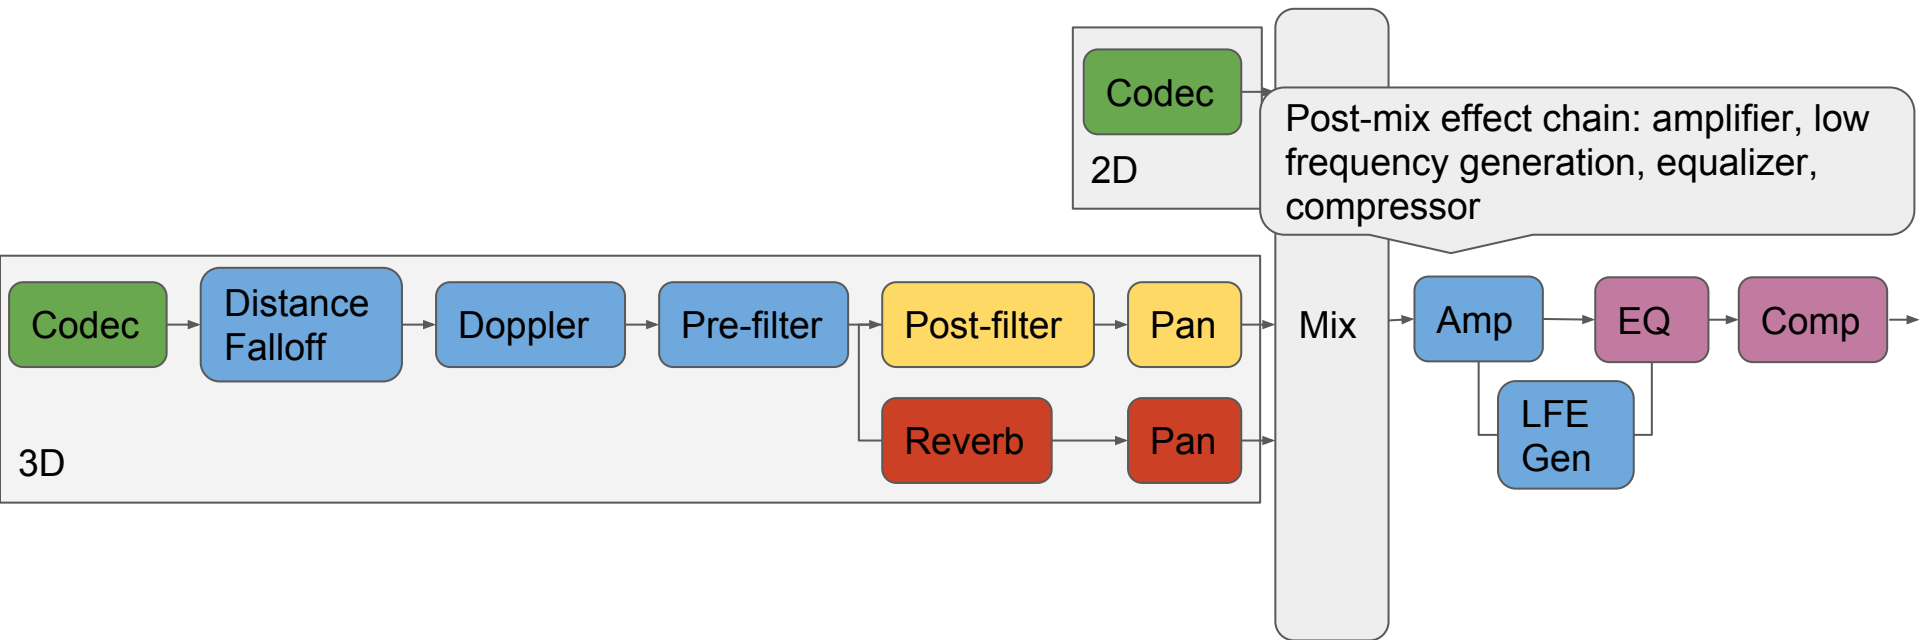
Sound engine architecture: Example Pipeline



Sound engine architecture: Example Pipeline



Sound engine architecture: Example Pipeline



Sound engine architecture: Notes

Some additional considerations:

- Buffering and latency
- Dynamic volume control
- Voice management
- Streaming

Sound engine architecture: Notes

Some additional considerations:

- **Buffering and latency**
- Dynamic volume control
- Voice management
- Streaming

Typically ring buffer stores final output:

- Push newly mixed audio on tail
- DAC consumes from head

Ring buffer size:

- Larger = more latency
- Smaller = greater chance of dropout

Sound engine architecture: Notes

Some additional considerations:

- Buffering and latency
- **Dynamic volume control**
- Voice management
- Streaming

Gameplay considerations:

- Hear an NPC whispering from far away
- Hear important cues over gun fire
- Quiet and loud sections

Must dynamically adjust volume:

- Duck environmental volume when important VO is happening
- Smoothly balance total sound pressure
- Optionally override distance attenuation

Sound engine architecture: Notes

Some additional considerations:

- Buffering and latency
- Dynamic volume control
- **Voice management**
- Streaming
- Hundreds / thousands of sound sources
- Can't afford to mix them all
- Prioritize sounds
- Playing new high priority sounds to steal voices from low priority sounds

Sound engine architecture: Notes

Some additional considerations:

- Buffering and latency
 - Dynamic volume control
 - Voice management
 - **Streaming**
- Audio is large
 - For large clips:
 - Don't have the whole thing resident
 - Stream into a ring buffer
 - But watch out for:
 - Bandwidth limitations
 - Seek times on physical disks

Ad hoc methods

Proper analytic models are expensive. Hack it.

- Sources that get too close to listeners
- Sounds from above and below
- Handling sound occluded by environment

Ad hoc methods

Proper analytic models are expensive. Hack it.

- **Sources that get too close to listeners**
- Sounds from above and below
- Handling sound occluded by environment

Far from listener:

- Source is a point and panning is simple

Close to listener:

- Listener becomes a volume
- Source becomes a volume
- Intersection of volumes defines how sound should be panned
- Quick and dirty approximations are fine

Ad hoc methods

Proper analytic models are expensive. Hack it.

- Sources that get too close to listeners
 - **Sounds from above and below**
 - Handling sound occluded by environment
- 7.1 surround is purely horizontal
 - Sound sources are projected onto horizontal plane
 - Post-filter & reverb clue listener into verticality

Ad hoc methods

Proper analytic models are expensive. Hack it.

- Sources that get too close to listeners
- Sounds from above and below
- **Handling sound occluded by environment**

Baseline:

- Tag reverb regions in game world

Dynamically modify filters based on:

- Look for LOS blockage between source and listener -- post-filter attenuates dry portion of sound
- Probabilistically model indirect occlusion with raycasts against collision world -- adjust reverb parameters on wet portion of sound

Summary

- Audio as signal in a linear time-invariant system
- Audio in 3D space
- Pseudo-API for a sound system
- Pseudo-pipeline for a sound system
- Convenient hacks and other considerations

End Lecture

<http://sol.gfxile.net/soloud/>